

Cloud Native Deployments of Cassandra using Kubernetes

Table of Contents

- [Prerequisites](#)
- [Cassandra Docker](#)
- [Quickstart](#)
- [Step 1: Create a Cassandra Headless Service](#)
- [Step 2: Use a StatefulSet to create Cassandra Ring](#)
- [Step 3: Validate and Modify The Cassandra StatefulSet](#)
- [Step 4: Delete Cassandra StatefulSet](#)
- [Step 5: Use a Replication Controller to create Cassandra node pods](#)
- [Step 6: Scale up the Cassandra cluster](#)
- [Step 7: Delete the Replication Controller](#)
- [Step 8: Use a DaemonSet instead of a Replication Controller](#)
- [Step 9: Resource Cleanup](#)
- [Seed Provider Source](#)

The following document describes the development of a *cloud native* [Cassandra](#) deployment on Kubernetes. When we say *cloud native*, we mean an application which understands that it is running within a cluster manager, and uses this cluster management infrastructure to

help implement the application. In particular, in this instance, a custom Cassandra `SeedProvider` is used to enable Cassandra to dynamically discover new Cassandra nodes as they join the cluster.

This example also uses some of the core components of Kubernetes:

- [*Pods*](#)
- [*Services*](#)
- [*Replication Controllers*](#)
- [*Stateful Sets*](#)
- [*Daemon Sets*](#)

Prerequisites

This example assumes that you have a Kubernetes version ≥ 1.2 cluster installed and running, and that you have installed the `kubectl` command line tool somewhere in your path. Please see the [getting started guides](#) for installation instructions for your platform.

This example also has a few code and configuration files needed. To avoid typing these out, you can `git clone` the Kubernetes repository to your local computer.

Cassandra Docker

The pods use the `gcr.io/google-samples/cassandra:v12` image from Google's [container registry](#).

The docker is based on `debian:jessie` and includes OpenJDK 8. This image

includes a standard Cassandra installation from the Apache Debian repo. Through the use of environment variables you are able to change values that are inserted into the `cassandra.yaml`.

ENV VAR	DEFAULT VALUE
CASSANDRA_CLUSTER_NAME	'Test Cluster'
CASSANDRA_NUM_TOKENS	32
CASSANDRA_RPC_ADDRESS	0.0.0.0

Quickstart

If you want to jump straight to the commands we will run, here are the steps:

```
#  
# StatefulSet  
#  
  
# create a service to track all cassandra statefulset nodes
```

```
kubectl create -f examples/storage/cassandra/cassandra-service.yaml
```

```
# create a statefulset
```

```
kubectl create -f examples/storage/cassandra/cassandra-statefulset.yaml
```

```
# validate the Cassandra cluster. Substitute the name of one of your pods.
```

```
kubectl exec -ti cassandra-0 -- nodetool status
```

```
# cleanup
```

```
grace=$(kubectl get po cassandra-0 --template '{{.spec.terminationGracePeriodSeconds}}') \  
  && kubectl delete statefulset,po -l app=cassandra \  
  && echo "Sleeping $grace" \  
  && sleep $grace \  
  && kubectl delete pvc -l app=cassandra
```

```
#
```

```
# Resource Controller Example
```

```
#
```

```
# create a replication controller to replicate cassandra nodes
```

```
kubectl create -f examples/storage/cassandra/cassandra-controller.yaml
```

```
# validate the Cassandra cluster. Substitute the name of  
one of your pods.
```

```
kubectl exec -ti cassandra-xxxxx -- nodetool status
```

```
# scale up the Cassandra cluster
```

```
kubectl scale rc cassandra --replicas=4
```

```
# delete the replication controller
```

```
kubectl delete rc cassandra
```

```
#
```

```
# Create a DaemonSet to place a cassandra node on each ku  
bernetes node
```

```
#
```

```
kubectl create -f examples/storage/cassandra/cassandra-da  
emonset.yaml --validate=false
```

```
# resource cleanup
```

```
kubectl delete service -l app=cassandra
```

```
kubectl delete daemonset cassandra
```

Step 1: Create a Cassandra Headless Service

A Kubernetes [*Service*](#) describes a set of [*Pods*](#) that perform the same task. In

Kubernetes, the atomic unit of an application is a Pod: one or more containers that *must* be scheduled onto the same host.

The Service is used for DNS lookups between Cassandra Pods, and Cassandra clients within the Kubernetes Cluster.

Here is the service description:

<!-- BEGIN MUNGE: EXAMPLE cassandra-service.yaml -->

```
apiVersion: v1
kind: Service
metadata:
  labels:
    app: cassandra
  name: cassandra
spec:
  clusterIP: None
  ports:
    - port: 9042
  selector:
    app: cassandra
```

[Download example](#)

<!-- END MUNGE: EXAMPLE cassandra-service.yaml -->

Create the service for the StatefulSet:

```
$ kubectl create -f examples/storage/cassandra/cassandra-service.yaml
```

The following command shows if the service has been created.

```
$ kubectl get svc cassandra
```

The response should be like:

NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
cassandra	None	<none>	9042/TCP	45s

If an error is returned the service create failed.

Step 2: Use a StatefulSet to create Cassandra Ring

StatefulSets (previously PetSets) are a feature that was upgraded to a *Beta* component in

Kubernetes 1.5. Deploying stateful distributed applications, like Cassandra, within a clustered

environment can be challenging. We implemented StatefulSet to greatly simplify this

process. Multiple StatefulSet features are used within this example,

but is out of

scope of this documentation. [Please refer to the Stateful Set documentation.](#)

The StatefulSet manifest that is included below, creates a Cassandra ring that consists of three pods.

This example includes using a GCE Storage Class, please update appropriately depending on the cloud you are working with.

<!-- BEGIN MUNGE: EXAMPLE cassandra-statefulset.yaml -->

```
apiVersion: "apps/v1beta1"
kind: StatefulSet
metadata:
  name: cassandra
spec:
  serviceName: cassandra
  replicas: 3
  template:
    metadata:
      labels:
        app: cassandra
    spec:
      containers:
        - name: cassandra
```



```
image: gcr.io/google-samples/cassandra:v12
imagePullPolicy: Always
ports:
  - containerPort: 7000
    name: intra-node
  - containerPort: 7001
    name: tls-intra-node
  - containerPort: 7199
    name: jmx
  - containerPort: 9042
    name: cql
resources:
  limits:
    cpu: "500m"
    memory: 1Gi
  requests:
    cpu: "500m"
    memory: 1Gi
securityContext:
  capabilities:
    add:
      - IPC_LOCK
lifecycle:
  preStop:
    exec:
      command: ["/bin/sh", "-c", "PID=$(pidof java) && kill $PID && while ps -p $PID > /dev/null; do sleep 1; done"]
```

env:

- name: MAX_HEAP_SIZE

value: 512M

- name: HEAP_NEWSIZE

value: 100M

- name: CASSANDRA_SEEDS

value: "cassandra-0.cassandra.default.svc.cluster.local"

- name: CASSANDRA_CLUSTER_NAME

value: "K8Demo"

- name: CASSANDRA_DC

value: "DC1-K8Demo"

- name: CASSANDRA_RACK

value: "Rack1-K8Demo"

- name: CASSANDRA_AUTO_BOOTSTRAP

value: "false"

- name: POD_IP

valueFrom:

fieldRef:

fieldPath: status.podIP

readinessProbe:

exec:

command:

- /bin/bash

- -c

- /ready-probe.sh

initialDelaySeconds: 15

timeoutSeconds: 5

```

    # These volume mounts are persistent. They are li
ke inline claims,
    # but not exactly because the names need to match
exactly one of
    # the stateful pod volumes.
    volumeMounts:
      - name: cassandra-data
        mountPath: /cassandra_data
    # These are converted to volume claims by the controlle
r
    # and mounted at the paths mentioned above.
    # do not use these in production until ssd GCEPersisten
tDisk or other ssd pd
    volumeClaimTemplates:
      - metadata:
          name: cassandra-data
          annotations:
            volume.beta.kubernetes.io/storage-class: fast
        spec:
          accessModes: [ "ReadWriteOnce" ]
          resources:
            requests:
              storage: 1Gi
    ---
kind: StorageClass
apiVersion: storage.k8s.io/v1beta1
metadata:
  name: fast

```

```
provisioner: kubernetes.io/gce-pd
parameters:
  type: pd-ssd
```

[Download example](#)

```
<!-- END MUNGE: EXAMPLE cassandra-statefulset.yaml -->
```

Create the Cassandra StatefulSet as follows:

```
$ kubectl create -f examples/storage/cassandra/cassandra-
statefulset.yaml
```

Step 3: Validate and Modify The Cassandra StatefulSet

Deploying this StatefulSet shows off two of the new features that StatefulSets provides.

1. The pod names are known
2. The pods deploy in incremental order

First validate that the StatefulSet has deployed, by running `kubectl` command below.

```
$ kubectl get statefulset cassandra
```

The command should respond like:

NAME	DESIRED	CURRENT	AGE
cassandra	3	3	13s

Next watch the Cassandra pods deploy, one after another. The StatefulSet resource deploys pods in a number fashion: 1, 2, 3, etc. If you execute the following command before the pods deploy you are able to see the ordered creation.

```
$ kubectl get pods -l="app=cassandra"
```

NAME	READY	STATUS	RESTARTS	AGE
cassandra-0	1/1	Running	0	1m
cassandra-1	0/1	ContainerCreating	0	8s

The above example shows two of the three pods in the Cassandra StatefulSet deployed.

Once all of the pods are deployed the same command will respond with the full StatefulSet.

```
$ kubectl get pods -l="app=cassandra"
```

NAME	READY	STATUS	RESTARTS	AGE
------	-------	--------	----------	-----

cassandra-0	1/1	Running	0	10m
cassandra-1	1/1	Running	0	9m
cassandra-2	1/1	Running	0	8m

Running the Cassandra utility `nodetool` will display the status of the ring.

```
$ kubectl exec cassandra-0 -- nodetool status
Datacenter: DC1-K8Demo
=====
Status=Up/Down
|/ State=Normal/Leaving/Joining/Moving
-- Address      Load          Tokens         Owns (effective)  Host ID                               Rack
UN  10.4.2.4      65.26 KiB     32             63.7%             9d27f81-6783-461d-8583-87de2589133e   Rack1-K8Demo
UN  10.4.0.4      102.04 KiB    32             66.7%             5559a58c-8b03-47ad-bc32-c621708dc2e4   Rack1-K8Demo
UN  10.4.1.4      83.06 KiB     32             69.6%             dce943c-581d-4c0e-9543-f519969cc805   Rack1-K8Demo
```

You can also run `cqlsh` to describe the keyspaces in the cluster.

```
$ kubectl exec cassandra-0 -- cqlsh -e 'desc keyspaces'

system_traces  system_schema  system_auth    system         system
_distributed
```

In order to increase or decrease the size of the Cassandra StatefulSet, you must use

`kubectl edit`. You can find more information about the edit command in the [documentation](#).

Use the following command to edit the StatefulSet.

```
$ kubectl edit statefulset cassandra
```

This will create an editor in your terminal. The line you are looking to change is

`replicas`. The example does not contain the entire contents of the terminal window, and the last line of the example below is the replicas line that you want to change.

```
# Please edit the object below. Lines beginning with a '#'
# will be ignored,
# and an empty file will abort the edit. If an error occurs
# while saving this file will be
# reopened with the relevant failures.
#
apiVersion: apps/v1beta1
kind: StatefulSet
metadata:
  creationTimestamp: 2016-08-13T18:40:58Z
```

```
generation: 1
labels:
  app: cassandra
name: cassandra
namespace: default
resourceVersion: "323"
selfLink: /apis/apps/v1beta1/namespaces/default/statefu
lsets/cassandra
uid: 7a219483-6185-11e6-a910-42010a8a0fc0
spec:
  replicas: 3
```

Modify the manifest to the following, and save the manifest.

```
spec:
  replicas: 4
```

The StatefulSet will now contain four pods.

```
$ kubectl get statefulset cassandra
```

The command should respond like:

NAME	DESIRED	CURRENT	AGE
cassandra	4	4	36m

For the Kubernetes 1.5 release, the beta StatefulSet resource does not have `kubectl scale` functionality, like a Deployment, ReplicaSet, Replication Controller, or Job.

Step 4: Delete Cassandra StatefulSet

Deleting and/or scaling a StatefulSet down will not delete the volumes associated with the StatefulSet. This is done to ensure safety first, your data is more valuable than an auto purge of all related StatefulSet resources. Deleting the Persistent Volume Claims may result in a deletion of the associated volumes, depending on the storage class and reclaim policy. You should never assume ability to access a volume after claim deletion.

Use the following commands to delete the StatefulSet.

```
$ grace=$(kubectl get po cassandra-0 --template '{{.spec.terminationGracePeriodSeconds}}') \
&& kubectl delete statefulset -l app=cassandra \
&& echo "Sleeping $grace" \
&& sleep $grace \
&& kubectl delete pvc -l app=cassandra
```

Step 5: Use a Replication

Controller to create Cassandra node pods

A Kubernetes

Replication Controller

is responsible for replicating sets of identical pods. Like a Service, it has a selector query which identifies the members of its set. Unlike a Service, it also has a desired number of replicas, and it will create or delete Pods to ensure that the number of Pods matches up with its desired state.

The Replication Controller, in conjunction with the Service we just defined, will let us easily build a replicated, scalable Cassandra cluster.

Let's create a replication controller with two initial replicas.

<!-- BEGIN MUNGE: EXAMPLE cassandra-controller.yaml -->

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: cassandra
  # The labels will be applied automatically
  # from the labels in the pod template, if not set
  # labels:
```

```
# app: cassandra

spec:
  replicas: 2
  # The selector will be applied automatically
  # from the labels in the pod template, if not set.
  # selector:
    # app: cassandra
  template:
    metadata:
      labels:
        app: cassandra
    spec:
      containers:
        - command:
            - /run.sh
          resources:
            limits:
              cpu: 0.5
          env:
            - name: MAX_HEAP_SIZE
              value: 512M
            - name: HEAP_NEWSIZE
              value: 100M
            - name: CASSANDRA_SEED_PROVIDER
              value: "io.k8s.cassandra.KubernetesSeedProvider"
            - name: POD_NAMESPACE
              valueFrom:
```

```

        fieldRef:
          fieldPath: metadata.namespace
      - name: POD_IP
        valueFrom:
          fieldRef:
            fieldPath: status.podIP
    image: gcr.io/google-samples/cassandra:v12
    name: cassandra
    ports:
      - containerPort: 7000
        name: intra-node
      - containerPort: 7001
        name: tls-intra-node
      - containerPort: 7199
        name: jmx
      - containerPort: 9042
        name: cql
    volumeMounts:
      - mountPath: /cassandra_data
        name: data
  volumes:
    - name: data
      emptyDir: {}

```

[Download example](#)

<!-- END MUNGE: EXAMPLE cassandra-controller.yaml -->

There are a few things to note in this description.

The `selector` attribute contains the controller's selector query. It can be explicitly specified, or applied automatically from the labels in the pod template if not set, as is done here.

The pod template's label, `app:cassandra`, matches the Service selector from Step 1. This is how pods created by this replication controller are picked up by the Service."

The `replicas` attribute specifies the desired number of replicas, in this case 2 initially. We'll scale up to more shortly.

Create the Replication Controller:

```
$ kubectl create -f examples/storage/cassandra/cassandra-controller.yaml
```

You can list the new controller:

```
$ kubectl get rc -o wide
```

NAME	DESIRED	CURRENT	AGE	CONTAINER(S)
IMAGE(S)				SELECTOR
cassandra	2	2	11s	cassandra
gcr.io/google-samples/cassandra:v12				app=cassandra

Now if you list the pods in your cluster, and filter to the label `app=cassandra`, you should see two Cassandra pods. (The `wide` argument lets you see which Kubernetes nodes the pods were scheduled onto.)

```
$ kubectl get pods -l="app=cassandra" -o wide
```

NAME	READY	STATUS	RESTARTS	AGE
NODE				
cassandra-21qyy	1/1	Running	0	1m
kubernetes-minion-b286				
cassandra-q6sz7	1/1	Running	0	1m
kubernetes-minion-9ye5				

Because these pods have the label `app=cassandra`, they map to the service we defined in Step 1.

You can check that the Pods are visible to the Service using the

following service endpoints query:

```
$ kubectl get endpoints cassandra -o yaml
apiVersion: v1
kind: Endpoints
metadata:
  creationTimestamp: 2015-06-21T22:34:12Z
  labels:
    app: cassandra
  name: cassandra
  namespace: default
  resourceVersion: "944373"
  selfLink: /api/v1/namespaces/default/endpoints/cassandr
a
  uid: a3d6c25f-1865-11e5-a34e-42010af01bcc
subsets:
- addresses:
  - ip: 10.244.3.15
    targetRef:
      kind: Pod
      name: cassandra
      namespace: default
      resourceVersion: "944372"
      uid: 9ef9895d-1865-11e5-a34e-42010af01bcc
  ports:
  - port: 9042
    protocol: TCP
```

To show that the `SeedProvider` logic is working as intended, you can use the

`nodetool` command to examine the status of the Cassandra cluster. To do this,

use the `kubectl exec` command, which lets you run `nodetool` in one of your

Cassandra pods. Again, substitute `cassandra-xxxxx` with the actual name of one

of your pods.

```
$ kubectl exec -ti cassandra-xxxxx -- nodetool status
```

```
Datacenter: datacenter1
```

```
=====
```

```
Status=Up/Down
```

```
|/ State=Normal/Leaving/Joining/Moving
```

```
-- Address          Load           Tokens   Owns (effective)  Host
```

```
  ID                                     Rack
```

```
UN  10.244.0.5    74.09 KB    256      100.0%           86fe
```

```
da0f-f070-4a5b-bda1-2eeb0ad08b77 rack1
```

```
UN  10.244.3.3    51.28 KB    256      100.0%           dafe
```

```
3154-1d67-42e1-ac1d-78e7e80dce2b rack1
```


Step 6: Scale up the Cassandra cluster

Now let's scale our Cassandra cluster to 4 pods. We do this by telling the Replication Controller that we now want 4 replicas.

```
$ kubectl scale rc cassandra --replicas=4
```

You can see the new pods listed:

```
$ kubectl get pods -l="app=cassandra" -o wide
```

NAME	READY	STATUS	RESTARTS	AGE
NODE				
cassandra-21qyy	1/1	Running	0	6m
kubernetes-minion-b286				
cassandra-81m2l	1/1	Running	0	47s
kubernetes-minion-b286				
cassandra-8qoyy	1/1	Running	0	47s
kubernetes-minion-9ye5				
cassandra-q6sz7	1/1	Running	0	6m
kubernetes-minion-9ye5				

In a few moments, you can examine the Cassandra cluster status again, and see that the new pods have been detected by the custom `SeedProvider`:

```
$ kubectl exec -ti cassandra-xxxxx -- nodetool status
Datacenter: datacenter1
=====
Status=Up/Down
|/ State=Normal/Leaving/Joining/Moving
--  Address          Load          Tokens   Owns (effective)  Host ID
--  --
UN  10.244.0.6         51.67 KB      256      48.9%             d07b
23a5-56a1-4b0b-952d-68ab95869163 rack1
UN  10.244.1.5         84.71 KB      256      50.7%             e060
df1f-faa2-470c-923d-ca049b0f3f38 rack1
UN  10.244.1.6         84.71 KB      256      47.0%             83ca
1580-4f3c-4ec5-9b38-75036b7a297f rack1
UN  10.244.0.5         68.2 KB       256      53.4%             72ca
27e2-c72c-402a-9313-1e4b61c2f839 rack1
```

Step 7: Delete the Replication Controller

Before you start Step 5, **delete the replication controller** you

created above:

```
$ kubectl delete rc cassandra
```

Step 8: Use a DaemonSet instead of a Replication Controller

In Kubernetes, a *Daemon Set* can distribute pods onto Kubernetes nodes, one-to-one. Like a *ReplicationController*, it has a selector query which identifies the members of its set. Unlike a *ReplicationController*, it has a node selector to limit which nodes are scheduled with the templated pods, and replicates not based on a set number of pods, but rather assigns a single pod to each targeted node.

An example use case: when deploying to the cloud, the expectation is that instances are ephemeral and might die at any time. Cassandra is built to replicate data across the cluster to facilitate data redundancy, so that in the case that an instance dies, the data stored on the instance does not,

and the

cluster can react by re-replicating the data to other running nodes.

DaemonSet is designed to place a single pod on each node in the Kubernetes

cluster. That will give us data redundancy. Let's create a

DaemonSet to start our storage cluster:

<!-- BEGIN MUNGE: EXAMPLE cassandra-daemonset.yaml -->

```
apiVersion: extensions/v1beta1
```

```
kind: DaemonSet
```

```
metadata:
```

```
  labels:
```

```
    name: cassandra
```

```
  name: cassandra
```

```
spec:
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        app: cassandra
```

```
    spec:
```

```
      # Filter to specific nodes:
```

```
      # nodeSelector:
```

```
      # app: cassandra
```

```
      containers:
```

```
        - command:
```

```
          - /run.sh
```

```
env:
  - name: MAX_HEAP_SIZE
    value: 512M
  - name: HEAP_NEWSIZE
    value: 100M
  - name: CASSANDRA_SEED_PROVIDER
    value: "io.k8s.cassandra.KubernetesSeedProvider"
  - name: POD_NAMESPACE
    valueFrom:
      fieldRef:
        fieldPath: metadata.namespace
  - name: POD_IP
    valueFrom:
      fieldRef:
        fieldPath: status.podIP
image: gcr.io/google-samples/cassandra:v12
name: cassandra
ports:
  - containerPort: 7000
    name: intra-node
  - containerPort: 7001
    name: tls-intra-node
  - containerPort: 7199
    name: jmx
  - containerPort: 9042
    name: cql
# If you need it it is going away in C* 4.0
```

```
      #- containerPort: 9160
      # name: thrift
    resources:
      requests:
        cpu: 0.5
      volumeMounts:
        - mountPath: /cassandra_data
          name: data
    volumes:
      - name: data
        emptyDir: {}
```

[Download example](#)

```
<!-- END MUNGE: EXAMPLE cassandra-daemonset.yaml -->
```

Most of this DaemonSet definition is identical to the ReplicationController definition above; it simply gives the daemon set a recipe to use when it creates new Cassandra pods, and targets all Cassandra nodes in the cluster.

Differentiating aspects are the `nodeSelector` attribute, which allows the DaemonSet to target a specific subset of nodes (you can label nodes just like other resources), and the lack of a `replicas` attribute due to the 1-to-1 node-

pod relationship.

Create this DaemonSet:

```
$ kubectl create -f examples/storage/cassandra/cassandra-  
daemonset.yaml
```

You may need to disable config file validation, like so:

```
$ kubectl create -f examples/storage/cassandra/cassandra-  
daemonset.yaml --validate=false
```

You can see the DaemonSet running:

```
$ kubectl get daemonset
```

NAME	DESIRED	CURRENT	NODE-SELECTOR
cassandra	3	3	<none>

Now, if you list the pods in your cluster, and filter to the label `app=cassandra`, you should see one (and only one) new cassandra pod

for each

node in your network.

```
$ kubectl get pods -l="app=cassandra" -o wide
```

NAME	READY	STATUS	RESTARTS	AGE
NODE				
cassandra-ico4r	1/1	Running	0	4s
kubernetes-minion-rpo1				
cassandra-kitfh	1/1	Running	0	1s
kubernetes-minion-9ye5				
cassandra-tzw89	1/1	Running	0	2s
kubernetes-minion-b286				

To prove that this all worked as intended, you can again use the

nodetool

command to examine the status of the cluster. To do this, use the

kubectl exec command to run **nodetool** in one of your newly-

launched cassandra pods.

```
$ kubectl exec -ti cassandra-xxxxx -- nodetool status
Datacenter: datacenter1
=====
Status=Up/Down
|/ State=Normal/Leaving/Joining/Moving
```


--	Address	Load	Tokens	Owns (effective)	Host
ID				Rack	
UN	10.244.0.5	74.09 KB	256	100.0%	86fe
	da0f-f070-4a5b-bda1-2eeb0ad08b77			rack1	
UN	10.244.4.2	32.45 KB	256	100.0%	0b1b
	e71a-6ffb-4895-ac3e-b9791299c141			rack1	
UN	10.244.3.3	51.28 KB	256	100.0%	dafe
	3154-1d67-42e1-ac1d-78e7e80dce2b			rack1	

Note: This example had you delete the cassandra Replication Controller before

you created the DaemonSet. This is because – to keep this example simple – the

RC and the DaemonSet are using the same `app=cassandra` label (so that their pods map to the service we created, and so that the SeedProvider can identify them).

If we didn't delete the RC first, the two resources would conflict with respect to how many pods they wanted to have running. If we wanted, we could support running both together by using additional labels and selectors.

Step 9: Resource Cleanup

When you are ready to take down your resources, do the following:

```
$ kubectl delete service -l app=cassandra  
$ kubectl delete daemonset cassandra
```

Custom Seed Provider

A custom `SeedProvider`

is included for running Cassandra on top of Kubernetes. Only when you deploy Cassandra

via a replication control or a daemonset, you will need to use the custom seed provider.

In Cassandra, a `SeedProvider` bootstraps the gossip protocol that Cassandra uses to find other

Cassandra nodes. Seed addresses are hosts deemed as contact points. Cassandra

instances use the seed list to find each other and learn the topology of the

ring. The `KubernetesSeedProvider`

discovers Cassandra seeds IP addresses via the Kubernetes API, those Cassandra

instances are defined within the Cassandra Service.

Refer to the custom seed provider [README](#) for further

`KubernetesSeedProvider` configurations. For this example you should not need

to customize the Seed Provider configurations.

See the [image](#) directory of this example for specifics on how the container docker image was built and what it contains.

You may also note that we are setting some Cassandra parameters (`MAX_HEAP_SIZE` and `HEAP_NEWSIZE`), and adding information about the [namespace](#).

We also tell Kubernetes that the container exposes both the `CQL` and `Thrift` API ports. Finally, we tell the cluster manager that we need 0.1 cpu (0.1 core).

```
<!-- BEGIN MUNGE: GENERATED_ANALYTICS -->
```

```
<!-- END MUNGE: GENERATED_ANALYTICS -->
```