

Lab: Control Groups (cgroups)

***Difficulty:** Intermediate*

***Time:** Approximately 15 minutes*

Control Groups (cgroups) are a feature of the Linux kernel that allow you to limit the access processes and containers have to system resources such as CPU, RAM, IOPS and network.

In this lab you will use **cgroups** to limit the resources available to Docker containers. You will see how to pin a container to specific CPU cores, limit the number of CPU shares a container has, as well as how to prevent a *fork bomb* from taking down a Docker Host.

You will complete the following steps as part of this lab.

- [Step 1 - cgroups and the Docker CLI](#)
- [Step 2 - Max-out two CPUs](#)
- [Step 3 - Set CPU affinity](#)
- [Step 4 - CPU share constraints](#)
- [Step 5 - Docker Compose and cgroups](#)
- [Step 6 - Preventing a fork bomb](#)

Prerequisites

You will need all of the following to complete this lab:

- A Linux-based Docker Host with at least two CPU cores
- Docker 1.11 or higher
- The `htop` tool installed on your Docker Host (`apt-get install htop` on Ubuntu hosts)
- Root access on the Docker Host

Step 1: cgroups and the Docker CLI

The `docker run` command provides many flags that allow you to apply cgroup limitations to new containers. The following flags are of note for this lab:

```
$ docker run --help
...
--cpu-shares                CPU shares (relative weight)
...
--cpuset-cpus               CPUs in which to allow execution (0-3, 0,1)
...
--pids-limit                Tune container pids limit (set -1 for unlimited)
```

For more information on cgroup related flags available to the `docker run` command, see the [docker run reference guide](#).

Step 2: Max-out two CPUs

In this step you'll start a new container that will max out two CPU cores. You will need `htop` installed on your Docker Host, and you will need to be running a Docker Host that has two or more CPU cores. This step will still work if your Docker Host only has a single core. However, some of the `htop` outputs will be slightly different.

1. If you haven't already done so, install `htop` on your Docker Host.

Ubuntu with `apt`.

```
$ sudo apt-get install htop
```

CentOS with `yum`.

```
$ sudo yum install htop
```

2. Clone the lab's GitHub repo locally on your Docker Host and change into the `cgroups/cpu-stress` directory.

```
$ sudo git clone https://github.com/docker/labs.git
Cloning into 'labs'...
remote: Counting objects: 1638, done.
```

```
Receiving objects: 100% (1638/1638), 8.84 MiB | 2.37 MiB/s, done.  
Resolving deltas: 22% (33/148) d 0 (delta 0), pack-reused 1638  
Resolving deltas: 100% (148/148), done.  
Checking connectivity... done.  
Checking out files: 100% (1389/1389), done.  
  
$ cd labs/security/cgroups/cpu-stress
```

3. Verify that the directory has a single `Dockerfile` and a single `docker-compose.yml` file.

```
$ ls -l  
-rw-r--r-- 1 root root 23 Jul 11 09:49 docker-compose.yml  
-rw-r--r-- 1 root root 85 Jul 11 09:49 Dockerfile
```

4. Inspect the contents of the `Dockerfile`.

```
$ cat Dockerfile  
FROM ubuntu:latest  
  
RUN apt-get update && apt-get install -y stress  
  
CMD stress -c 2
```

As you can see, the Dockerfile describes a simple Ubuntu-based container that runs a single command `stress -c 2`. This command spawns two processes - both spinning on `sqrt()`. The effect of this command is to stress two CPU cores.

5. Build the image specified in the `Dockerfile`.

```
$ sudo docker build -t cpu-stress .  
Sending build context to Docker daemon 3.072 kB  
Step 1 : FROM ubuntu:latest  
latest: Pulling from library/ubuntu  
90d6565b970a: Pull complete  
40553bdb8474: Pull complete  
<snip>  
Step 3 : CMD stress -c 2  
---> Running in b90defcccb8  
---> 9e6a3f316e91  
Removing intermediate container b90defcccb8  
Successfully built 9e6a3f316e91
```

6. Run a new container called **stresser** based on the image built in the previous step.

```
$ sudo docker run -d --name stresser cpu-stress  
stress: info: [5] dispatching hogs: 2 cpu, 0 io, 0  
vm, 0 hdd
```

Be sure to run the container in the background with the `-d` flag so that you can run the `htop` command in the next step from the same terminal.

7. View the impact of the container using the `htop` command.

```
1 [|||||100.0%]
2 [0.0%]
3 [0.5%]
4 [|||||100.0%]
Mem[204/15041MB]
Swp[0/0MB]
Tasks: 37, 38 thr; 3 running
Load average: 0.98 0.53 0.27
Uptime: 00:07:20
```

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
9315	root	20	0	7484	96	0	R	100.	0.0	0:35.48	stress -c 2
9316	root	20	0	7484	96	0	R	100.	0.0	0:35.48	stress -c 2
9576	root	20	0	25208	4220	3000	R	0.5	0.0	0:00.08	htop

The output above shows two stress processes (**stress -c 2**) maxing out two of the CPUs on the system (CPU 1 and CPU 4). Both `stress` processes are in the running state, and both consuming 100% of the CPU they are executing on.

8. Stop and remove the **stresser** container.

```
$ sudo docker stop stresser && sudo docker rm stresser
stresser
stresser
```

You have seen how it is possible for a single container to max out CPU resources on a Docker Host. You would max out your entire Docker Host if you were to start a **stress** worker processes for each CPU core.

Step 3: Set CPU affinity

Docker makes it possible to restrict containers to a particular CPU core, or set of CPU cores. In this step you'll see how to restrict a container to a single CPU core using `docker run` with the `--cpuset-cpus` flag.

1. Run a new Docker container called **stresser** and restrict it to running on the first CPU on the system.

```
$ sudo docker run -d --name stresser --cpuset-cpus
0 cpu-stress

0bfbf2d33516065bbcf56bd8f9df24749312460141bca729f5
3d66a9b2dba6b
```

The `--cpuset-cpus` flag indexes CPU cores starting at 0. Therefore, CPU 0 is the first CPU on the system. You can specify multiple CPU cores as `0-4`, `0,3` etc.

2. Run the `htop` command to see the impact the container is having on the Docker Host.

```
1 [|||||||||||||||||||||||||||||||||||||100.0%] Tasks: 37, 38 thr; 3 running
2 [|||] 0.9%] Load average: 1.02 1.41 0.91
3 [ 0.0%] Uptime: 00:16:12
4 [ 0.0%]
Mem[|||||] 192/15041MB
Swp[ 0/0MB]

PID USER PRI NI VIRT RES SHR S CPU% MEM% TIME+ Command
9713 root 20 0 7484 96 0 R 50.2 0.0 0:06.41 stress -c 2
9712 root 20 0 7484 96 0 R 49.8 0.0 0:06.41 stress -c 2
9715 root 20 0 24796 3796 2976 R 0.5 0.0 0:00.04 htop
```

There are a few things worth noting about what you have just done:

- The container is based on the same **cpu-stress** image that spawns two stress worker processes.
- The two stress worker processes have been restricted to running against a single CPU core by the `--cpuset-cpus` flag.
- Each of the two stress processes is consuming ~50% of available time on the single CPU core they are executing on.
- `htop` indexes CPU cores starting at 1 whereas `--cpuset-cpus` indexes starting at 0.

You have seen how to lock a container to a executing on a single CPU core on the Docker Host. Feel free to experiment further with the `--cpuset-cpus` flag.

 Step 4: Set CPU share constraints

By default, all containers get an equal share of time executing on the Docker Host's CPUs. This allocation of time can be modified by

changing the container's CPU share weighting relative to the weighting of all other running containers.

The `--cpu-shares` flag takes a value from 0-1024. The default value is 1024, and a value of 0 will also default to 1024. If three containers are running and one has 1024 shares, while the other two have 512, the first container will get 50% of processor time while the other two will each get 25%. However, these shares are only enforced when CPU intensive tasks are running. When a container is not busy, other containers can use the free CPU time.

Shares of CPU time are balanced across all cores on a multi-core Docker Host.

In this step you will use the `docker run` command with the `--cpu-shares` flag to influence the amount of CPU time containers get. You will start one container (**container-1**) with 768 shares, and another container (**container-2**) with 256 shares. Each container will be locked to the first CPU on the system, and each will be running the `stress -c 2` process. **Container-1** will receive 75% of the CPU time whereas **container-2** will receive 25%.

1. Start the first container with 768 CPU shares.

```
$ sudo docker run -d --name container-1 --cpuset-cpus 0 --cpu-shares 768 cpu-stress
```

2. Start the second container with 256 CPU shares.

```
$ sudo docker run -d --name container-2 --cpuset-cpus 0 --cpu-shares 256 cpu-stress
```

3. Verify that both containers are running with the `docker ps` command.

```
$ sudo docker ps
```

CONTAINER ID	IMAGE	COMMAND
CREATED	STATUS	
PORTS	NAMES	

725dc16fac5a cpu-stress "/bin/sh -c 'stress -" 2 minutes ago Up 2 minutes container-2

f82f95757d3f cpu-stress "/bin/sh -c 'stress -" 2 minutes ago Up 2 minutes container-1

4. View the output of ``htop``.

Notice two things about the ``htop`` output. First, only a single CPU is being maxed out. Second, there are four ``st`

ress` processes running. The first two in the list equate to ~75% of CPU time, and the second two equate to ~25% of CPU time.

In this step you've seen how to use the `--cpu-shares` flag to influence the amount of CPU time containers get relative to other containers on the same Docker Host.

Feel free to experiment further by running more containers with different relative weights.

[Step 5: Docker Compose and cgroups](#)

In this step you'll see how to set CPU affinities via Docker Compose files.

See [this section](https://docs.docker.com/compose/compose-file/#cpu-shares-cpu-quota-cpuset-domainname-hostname-ipc-mac-address-mem-limit-memswap-limit-privileged-read-only-restart-shm-size-stdin-open-tty-user-working-dir) of the Docker Compose documentation for more information on leveraging other cgroup capabilities using Docker Compose.

Make sure you are in the `dockercon-workshop/cgroups/cpu-stress` directory of the repo that you pulled in Step 2.

1. Edit the `docker-compose.yml` file and add the `cpuset

```
: '3'` line as shown below:
```

```
cpu-stress:
```

```
build: .
```

```
cpuset: '3'
```

The above `docker-compose.yml` file will ensure that containers based from it will run on CPU core #3. You will obviously need a Docker Host with at least 4 CPU cores for this to work.

2. Bring the application up in the background.

```
```
```

```
$ sudo docker-compose up -d
```

```
Creating cpustress_cpu-stress_1
```

```
```
```

3. Run `htop` to see the effect of the `cpuset` parameter

```
.
```

```

```

The `htop` output above shows the container and it's two

``stress`` processes locked to CPU core 4 (``cpuset`` in Docker Compose indexes CPU cores starting at 0 whereas ``htop`` indexes CPU cores starting at 1).

In this step you've seen how Docker Compose can set container CPU affinities. Remember that Docker Compose can also set CPU quotas and shares. See the [documentation](<https://docs.docker.com/compose/compose-file/#cpu-shares-cpu-quota-cpuset-domainname-hostname-ipc-mac-address-mem-limit-memswap-limit-privileged-read-only-restart-shm-size-stdin-open-tty-user-working-dir>) for more detail.

``Step 6: Preventing a fork bomb

A *fork bomb* is a form of denial of service (DoS) attack where a process continually replicates itself with the goal of depleting system resources to the point where a system can no longer function.

In this step you will use the ``--pids-limit`` flag to limit the number of processes a container can fork at runtime. This will prevent a fork bomb from consuming the Docker Host's entire process table.

1. Start a new container and limit the number of processes it can create to 200 with the following command.

```
$ sudo docker run --rm -it --pids-limit 200 debian:jessie bash
```

Unable to find image 'debian:jessie' locally

jessie: Pulling from library/debian

5c90d4a2d1a8: Pull complete

Digest:

sha256:8b1fc3a7a55c42e3445155b2f8f40c55de5f8bc8012992b26b57053

Status: Downloaded newer image for debian:jessie

```
root@c0eb76d2481c:/#
```

****Do not proceed to the next step**** if you receive the following warning about pids limit support - **"WARNING: Your kernel does not support pids limit capabilities, pids limit discarded"**. If you receive this warning, you can [watch the demo](<https://asciinema.org/a/dewdpjzlzom4zasdz0c0c46jt9>) as a video instead.

2. Run the following command to create a fork bomb in the container you just started.

The previous command will have attached your terminal to the `bash` shell inside of the container created in the previous step.

```
**Do not complete this step if you received the error that your kernel does not support the pids limit feature...*  
*
```

```
root@c0eb76d2481c:/# :(){ :|: & };;
```

```
[1] 6
```

```
root@a3eeb655f301:/# bash: fork: retry: No child processes
```

```
bash: fork: retry: No child processes
```

```
bash: fork: retry: No child processes
```

```
bash: fork: retry: No child processes
```

```
bash: fork: retry: No child processes
```

```
bash: fork: retry: Resource temporarily unavailable
```

```
bash: fork: retry: No child processes
```

```
<snip>
```

```
bash: fork: retry: No child processes
```

```
bash: fork: retry: Resource temporarily unavailable
```

```
bash: fork: retry: No child processes
```

```
^C
```

```
[1]+ Done : | :
```

```
> **Note:** The command above is `:(){ :|: & };;`.
```

```
You will need to press `Ctrl-C` to stop the process.
```

In this step you have seen how to set a process limits on a container that will prevent it from consuming all process table resources on the underlying Docker Host.

Summary

Congratulations. You've seen how to use some of the `cgroups` features supported by Docker that allow you to limit and constrain a container's access to Docker Host system resources.