

Docker

A FRAMEWORK FOR DATA INTENSIVE COMPUTING

Agenda


Intro / Prep Environments

Day 1: Docker Deep Dive

Day 2: Kubernetes Deep Dive

Day 3: Advanced Kubernetes: Concepts, Management, Middleware

Day 4: Advanced Kubernetes: CI/CD, open discussions




What is This Docker Thing?

- ☐ A company?
- ☐ A format?
- ☐ An API?
- ☐ Framework?
- ☐ Or, maybe something in the Cloud.
- ☐ Everything's in the cloud **seems** like these days!

That's **Not** New About Docker

- ❑ Linux-**native** functionality
- ❑ Control Groups (**cgroups**)
- ❑ Kernel **Namespaces**
- ❑ **chroot**
- ❑ Linux **Capabilities**
- ❑ Security (**SELinux**)



What Are Cgroups

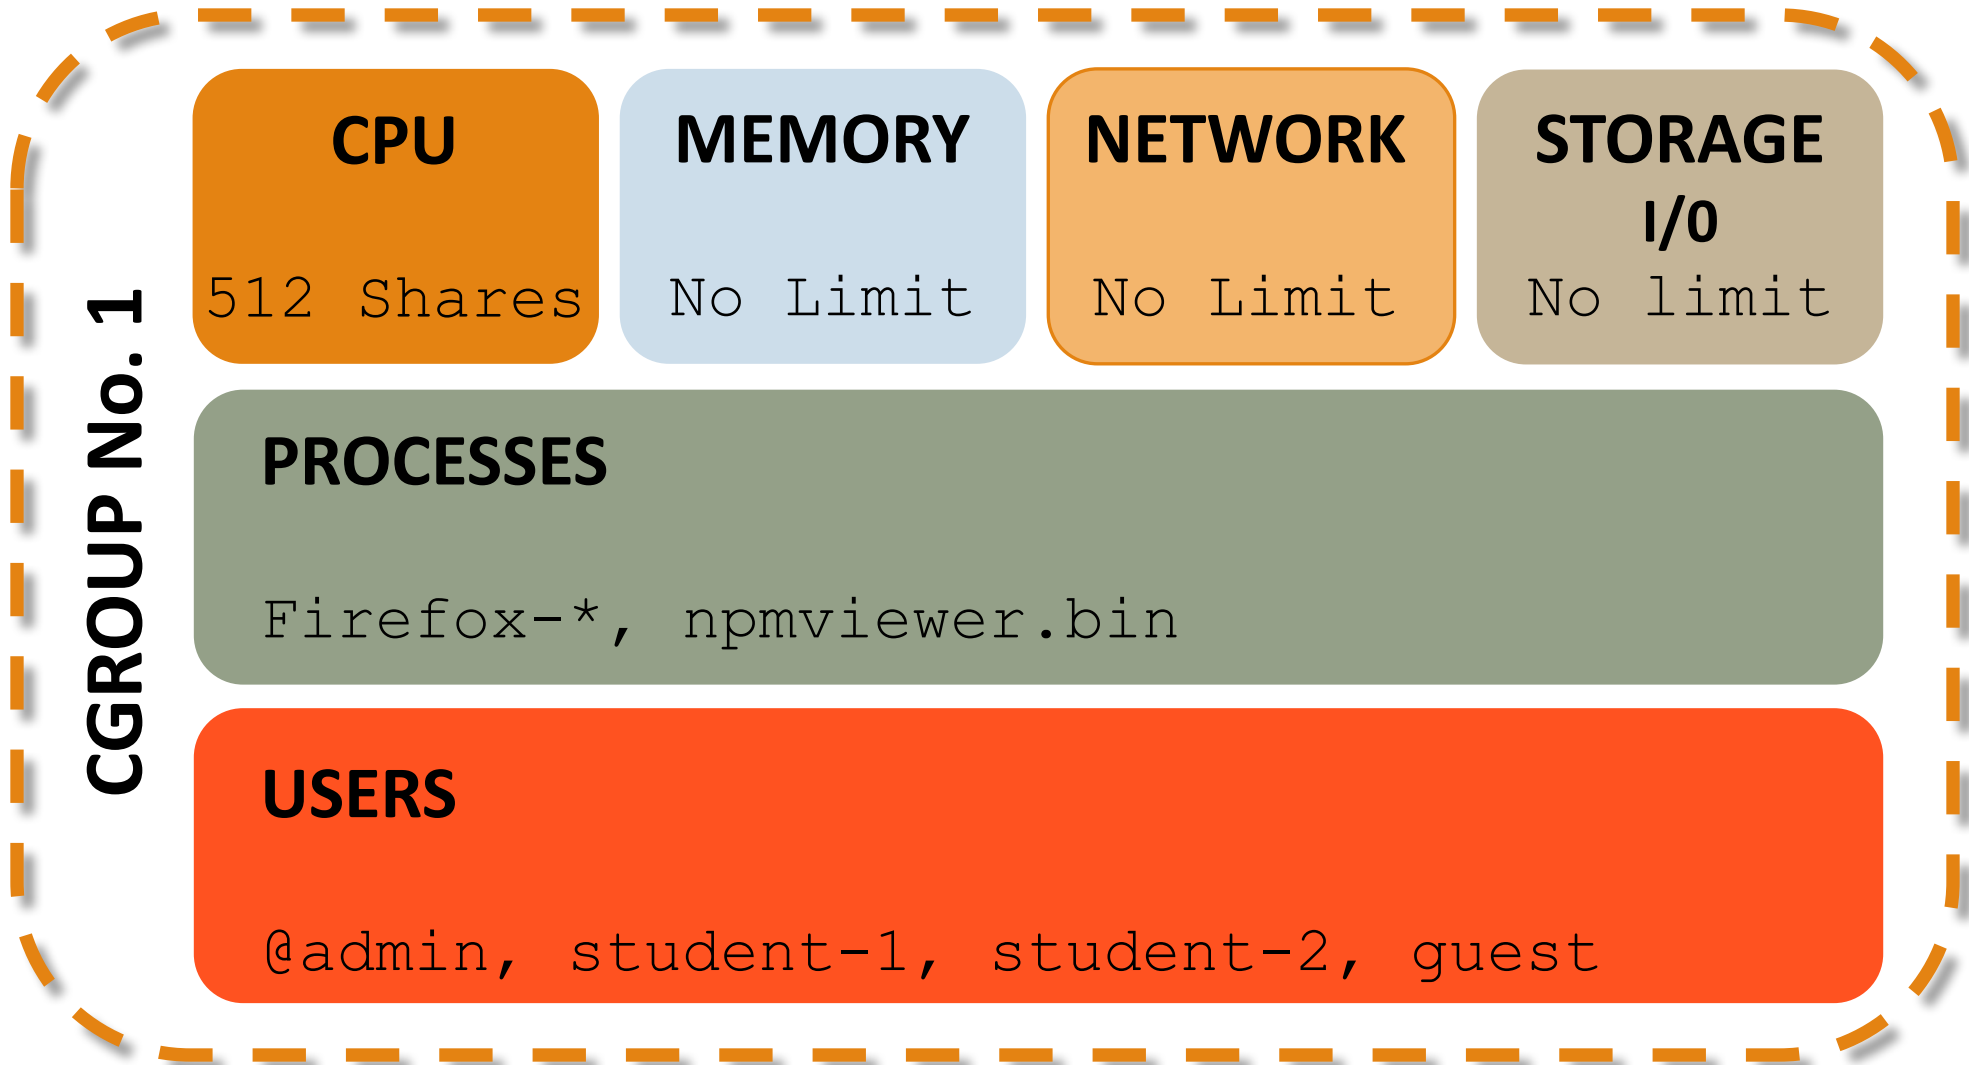
- ❑ Built into Kernel (RHEL7/Debian/etc)
- ❑ Generically isolates resource usage (CPU, memory, disk, network)
- ❑ Guarantee resources to app/set of apps
- ❑ Can be adjusted on the fly
- ❑ Can monitor the cgroup itself to see utilization

What Are Cgroups

□ Wikipedia says:

“**cgroups** (abbreviated from control groups) is a Linux kernel feature that limits, accounts for, and isolates the resource usage (CPU, memory, disk I/O, network, etc.) of a collection of processes.”

Diagram of Control Groups



Kernel Namespaces

- ❑ Isolating views of the system
- ❑ You can make a process think it's the only process
- ❑ Built-in ways to "virtualize" a process

Wikipedia says:

“**Namespaces** are a feature of the Linux **kernel** that isolates and virtualizes system resources of a collection of processes. Examples of resources that can be virtualized include process IDs, hostnames, user IDs, network access, inter-process communication, and filesystems.”

Kernel Namespaces

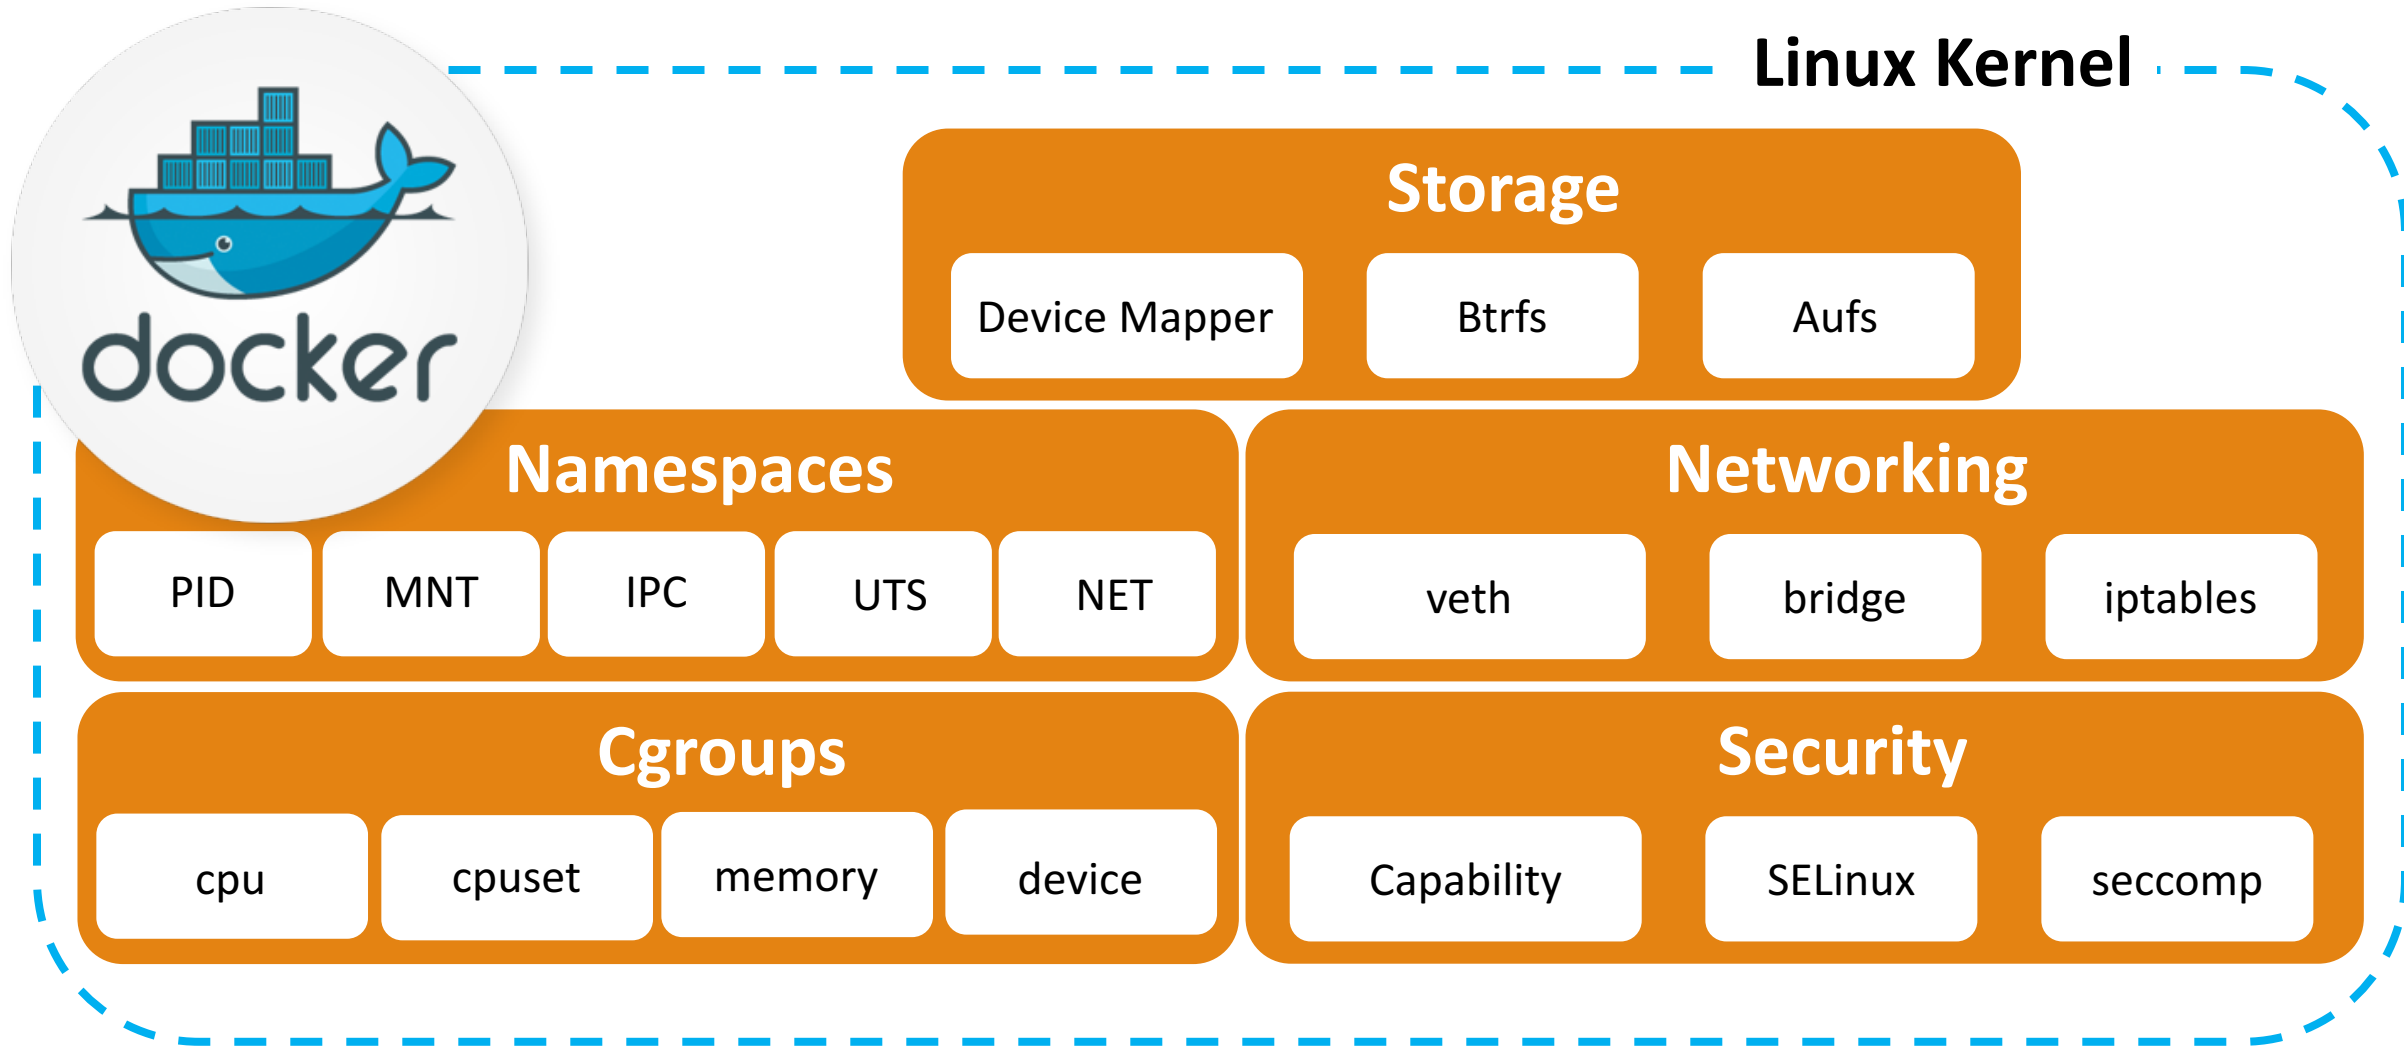
- ❑ **mnt** mount points, filesystem
- ❑ **pid** processes
- ❑ **net** network stack
- ❑ **ipc** inter-process comms
- ❑ **uts** hostname
- ❑ **user** UIDs



Linux capabilities

- ❑ Like Linux, `root` has **all** capabilities.
- ❑ Includes a fine-grained division of `root`'s permissions for a processes.
- ❑ `CAP_NET_ADMIN` - modify routing tables, firewalling, NAT, etc.
- ❑ `CAP_KILL` - bypass any checks for sending the kill signals
- ❑ `CAP_SYS_ADMIN` - mount, set hostname, etc.

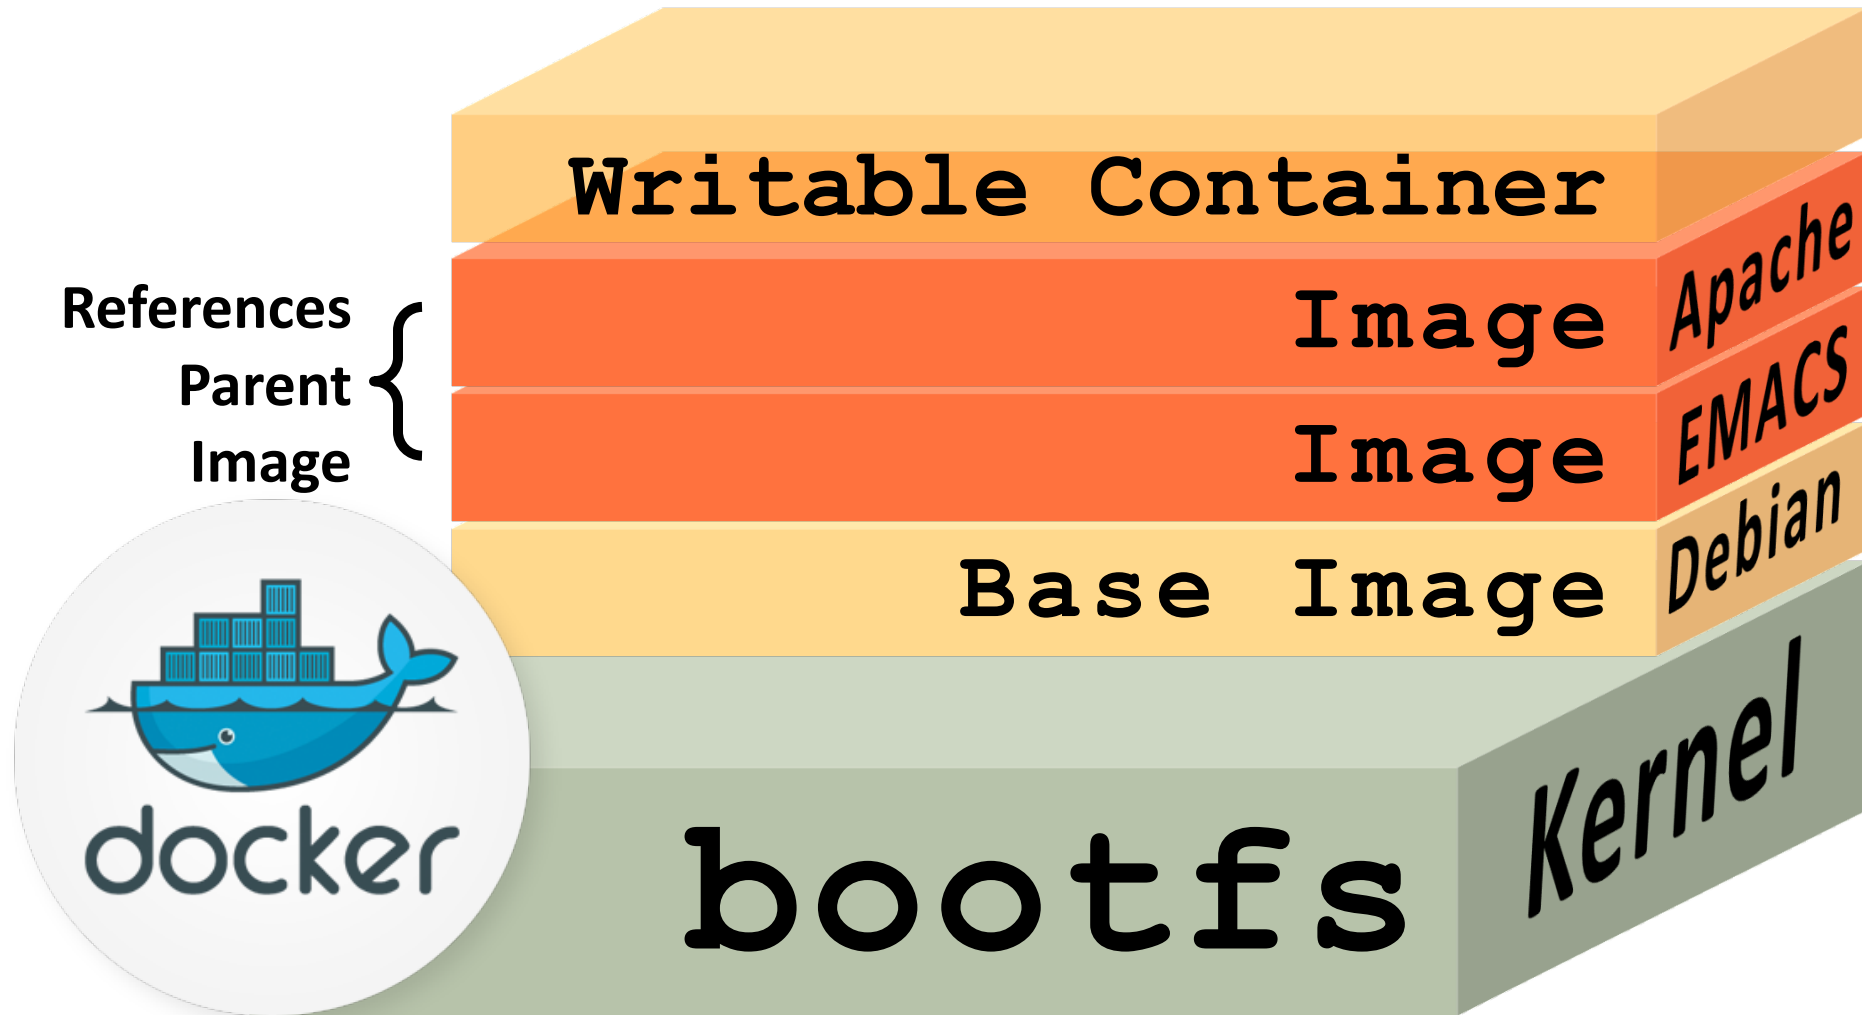
Docker Brings Together




Why is this Important?

- ❑ Think about image formatting vs. golden imaging
- ❑ When using API calls
- ❑ Packaging
- ❑ Structural separation of concerns (Devs/Ops)
- ❑ Density and infrastructure utilization

Foundations of Docker



Virtualization vs. Containerization



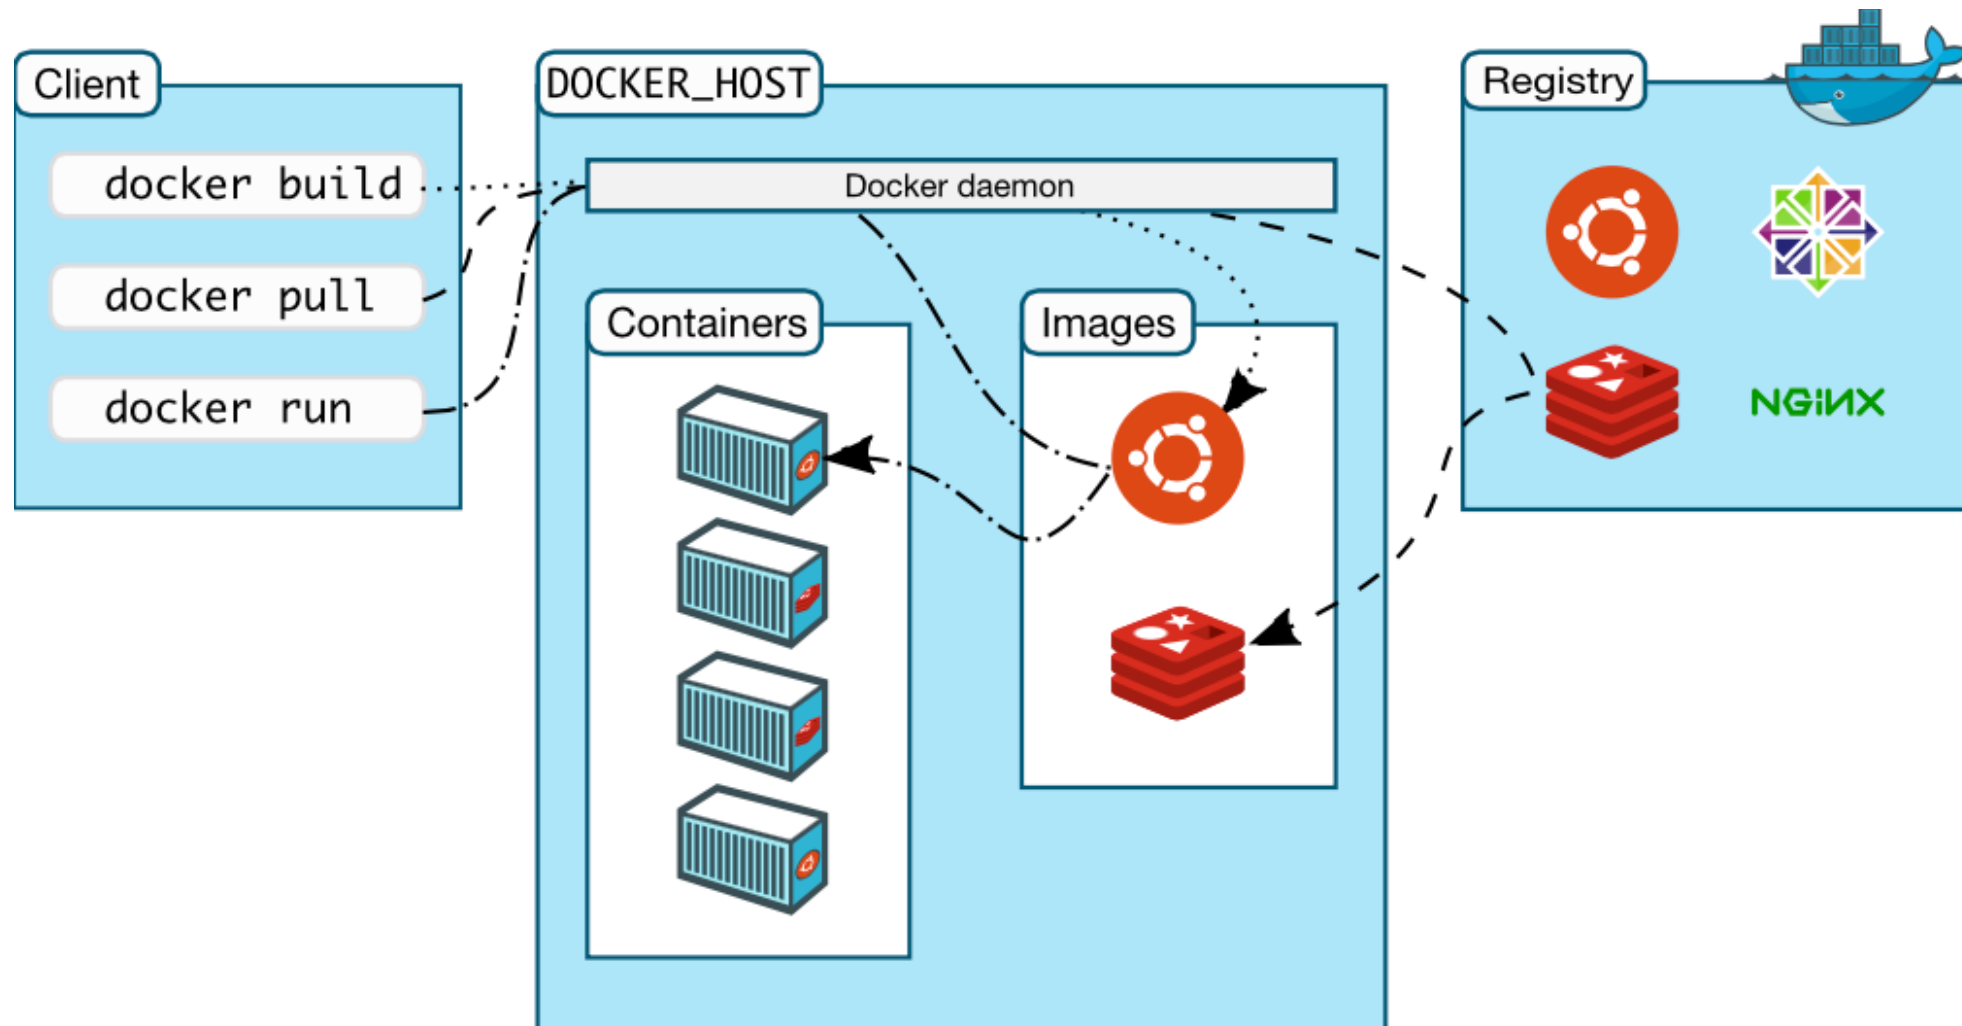
Immutable Infrastructure

- ❑ "We'll put it back in **Ansible**"
- ❑ Cattle **vs.** Pets
- ❑ Don't change it; **replace it**
- ❑ System created fully from automation; **avoid drift**
- ❑ Manual **intervention** is error prone
- ❑ How does **Docker** help?

Basic Docker Components

- ❑ Docker Client
- ❑ Docker Daemon
- ❑ Images
- ❑ Registry
- ❑ Containers

Basic Docker Components




Docker Images

- ❑ **Templates** from which containers are created
- ❑ **Layered** using union filesystems
- ❑ Each change to the system is a **layer** added
- ❑ Typically created from **instructions** stored in Dockerfiles
- ❑ Stored in a Docker **registry**, locally or remotely

Docker Containers

- ❑ **Runtime** instance of a Docker Image
- ❑ **Copy** on write file system; changes localized
- ❑ “**Virtualized**” with namespaces, cgroups, SELinux, etc.
- ❑ Each has **own** IP address, networking, and volumes
- ❑ Intended to run **single** process (process virtualization)



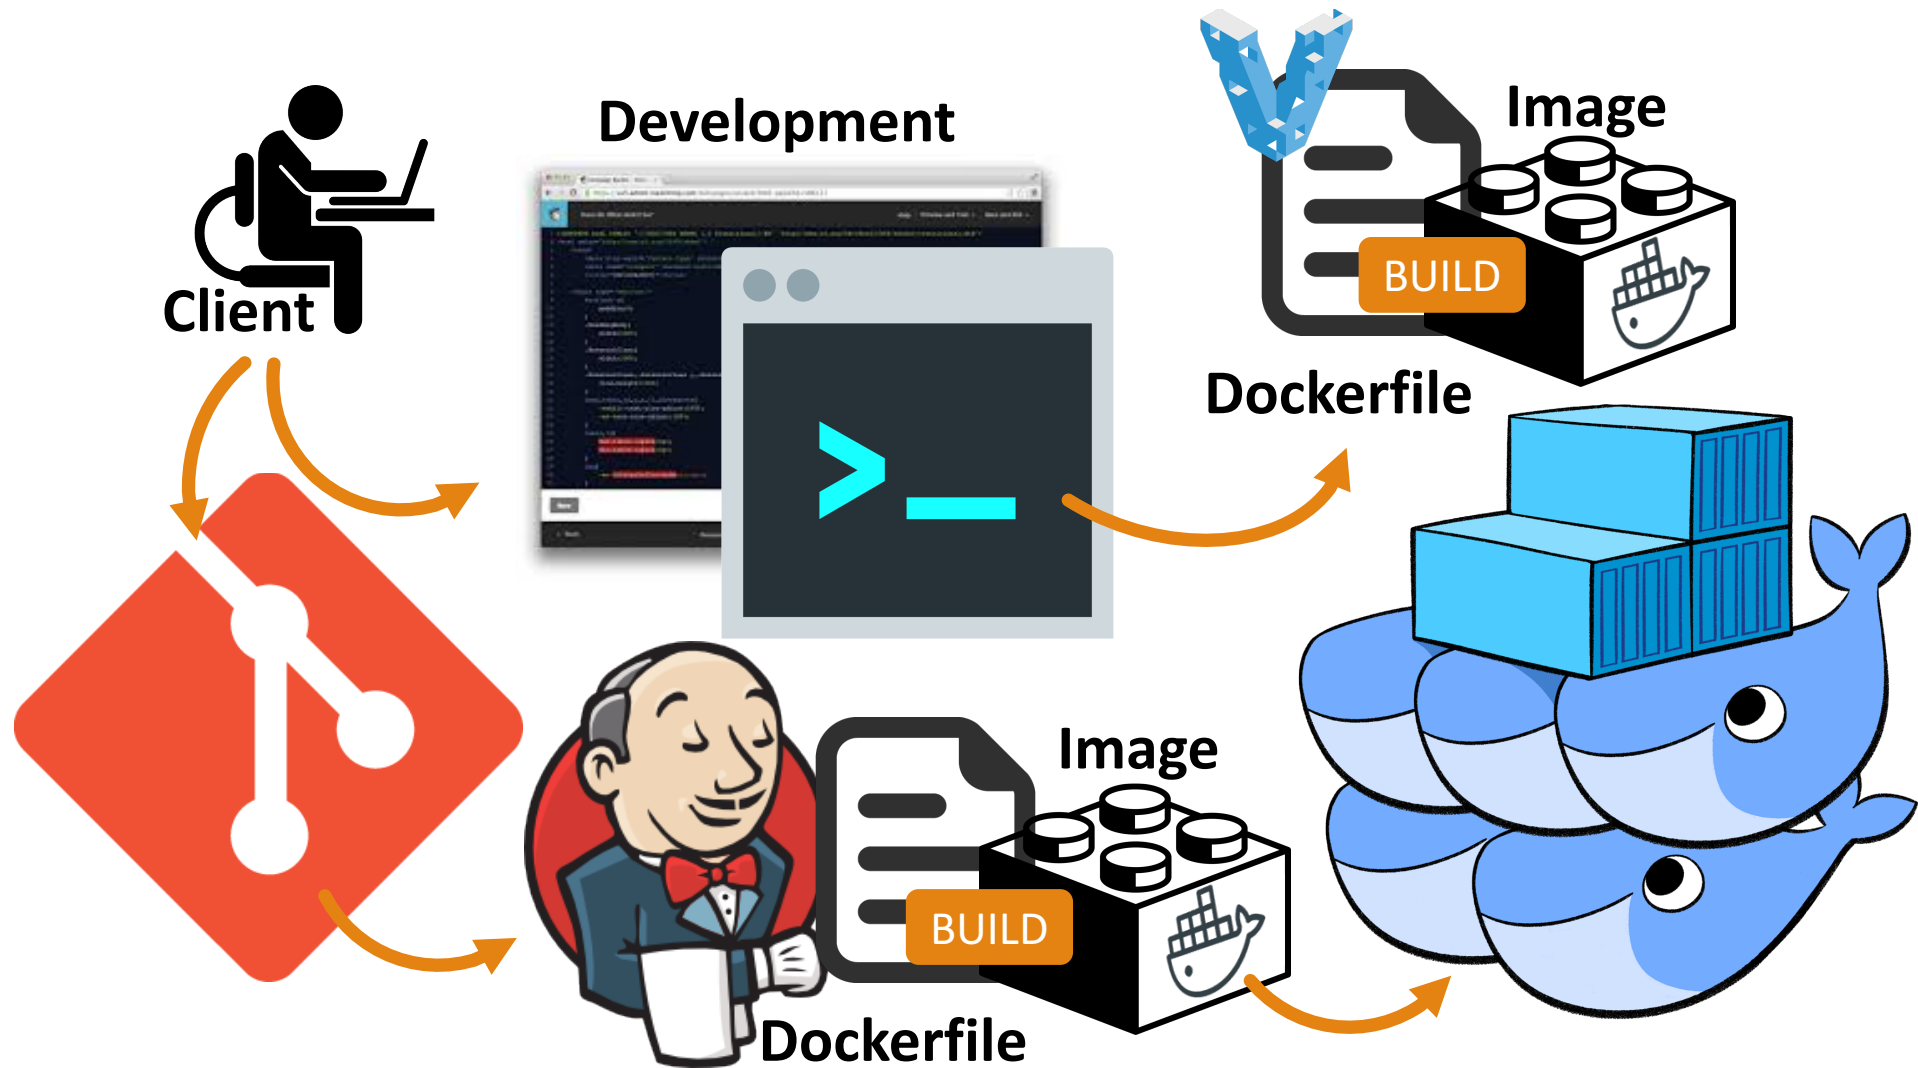
Developer Workflow

- ❑ Work from **Vagrant** image
- ❑ Can trash and reboot it any time
- ❑ **Locally** running Docker client
- ❑ Source code in developer IDE
- ❑ When ready, use tooling to **generate** Docker image, or hand craft
- ❑ Run image locally - possibly with others
- ❑ Push code, or image
- ❑ CI process kicks in

Developer Works Locally



Developer Pushes Code



Hello Docker Container

LET'S CREATE OUR FIRST CONTAINER


Pull a Docker Image

- ❑ Let's pull a Centos 7 image from DockerHub, like this:

```
docker pull centos:7
```

- ❑ Output:

```
Pulling from library/centos
fa5be2806d4c: Pull complete
0cd86ce0a197: Pull complete
e9407f1d4b65: Pull complete
c9853740aa05: Pull complete
e9fa5d3a0d0e: Pull complete
Digest: sha256:xxxxxxxxxxxxxxxx...
Status: Downloaded newer image for centos:7
```



List Docker Images

- Now, let's list all our locally installed images, like:

```
docker images
```

- Output:

REPOSITORY	TAG	IMAGE ID	CREATED	VIRTUAL SIZE
centos	7	e9fa5d3a0d0e	2 days ago	172.3 MB

List Docker Images

- ❑ We can pass the images command (-a) option to see all images, past and present:

```
docker images -a
```

- ❑ Output:

REPOSITORY	TAG	IMAGE ID	CREATED	VIRTUAL SIZE
centos	7	e9fa5d3a0d0e	2 days ago	172.3 MB
<none>	<none>	c9853740aa05	2 days ago	172.3 MB
<none>	<none>	e9407f1d4b65	2 days ago	172.3 MB
<none>	<none>	0cd86ce0a197	2 days ago	172.3 MB
<none>	<none>	fa5be2806d4c	5 weeks ago	0 B

Running Our First Docker Container

- Here's how we can then run a Linux command inside a Docker container:

```
docker run -rm centos:7 echo "hello world"
```

- Output:

```
hello world
```

- Woah, what happened? It just printed "hello, world"? So what?

Running Our First Docker Container

- ❑ Let's run a Linux shell **inside** a Docker container, from the inside:

```
docker run --it --rm centos:7 bash
```

- ❑ Output:

```
[root@d7dfcc490cbe /]# _
```

Inside Our Container

- Now from inside the Docker container itself, we'll issue commands, like this:

```
# ll /etc/*-release
```

- Output:

```
-rw-r--r-- 1 root root 38 Mar 31 2015 /etc/centos-release
-rw-r--r-- 1 root root 393 Mar 31 2015 /etc/os-release
lrwxrwxrwx 1 root root 14 Aug 14 21:00 /etc/redhat-release ->
centos-release
lrwxrwxrwx 1 root root 14 Aug 14 21:00 /etc/system-release ->
centos-release
```

Inside Our Container

- We can now play around with Linux commands within the container:

```
# hostname -f  
# cat /etc/hosts  
# ps aux  
# yum -y install vim  
# ip a
```

- If it looks like Linux and it works like Linux...

Let's Make a Mess

- Now, lets pretend we do some destructive stuff, on purpose:

```
# rm -fr /usr/sbin
```

- And now say we delete this, too:

```
# rm -fr /usr/bin
```

- Now, if we tried to issue any commands, like `ls`, `ps`, or `cd` we'd get:

```
bash: /usr/bin/ls: No such file or directory
```

- Whoops... cannot `ls`, `cd` or do anything useful anymore. What'd we'd do?

How Easy Is Starting Over

- ❑ In the last slide we pretended to destroy all our hard work.
- ❑ Now, how easy is it to start over?
- ❑ Just exit the container and fire up a new one:

```
docker run --it --rm centos:7 bash
```

- ❑ And everything is back! Wow... right?

Lab

End of Chapter
