

Docker Swarm

CONNECTING AND MANAGING MULTIPLE CONTAINERS

Agenda

Intro / Prep Environments

Day 1: Docker Deep Dive

Day 2: Docker Advanced Deep Dive

Day 3: Kubernetes Deep Dive

Day 4: Advanced Kubernetes: Concepts, Management, Middleware

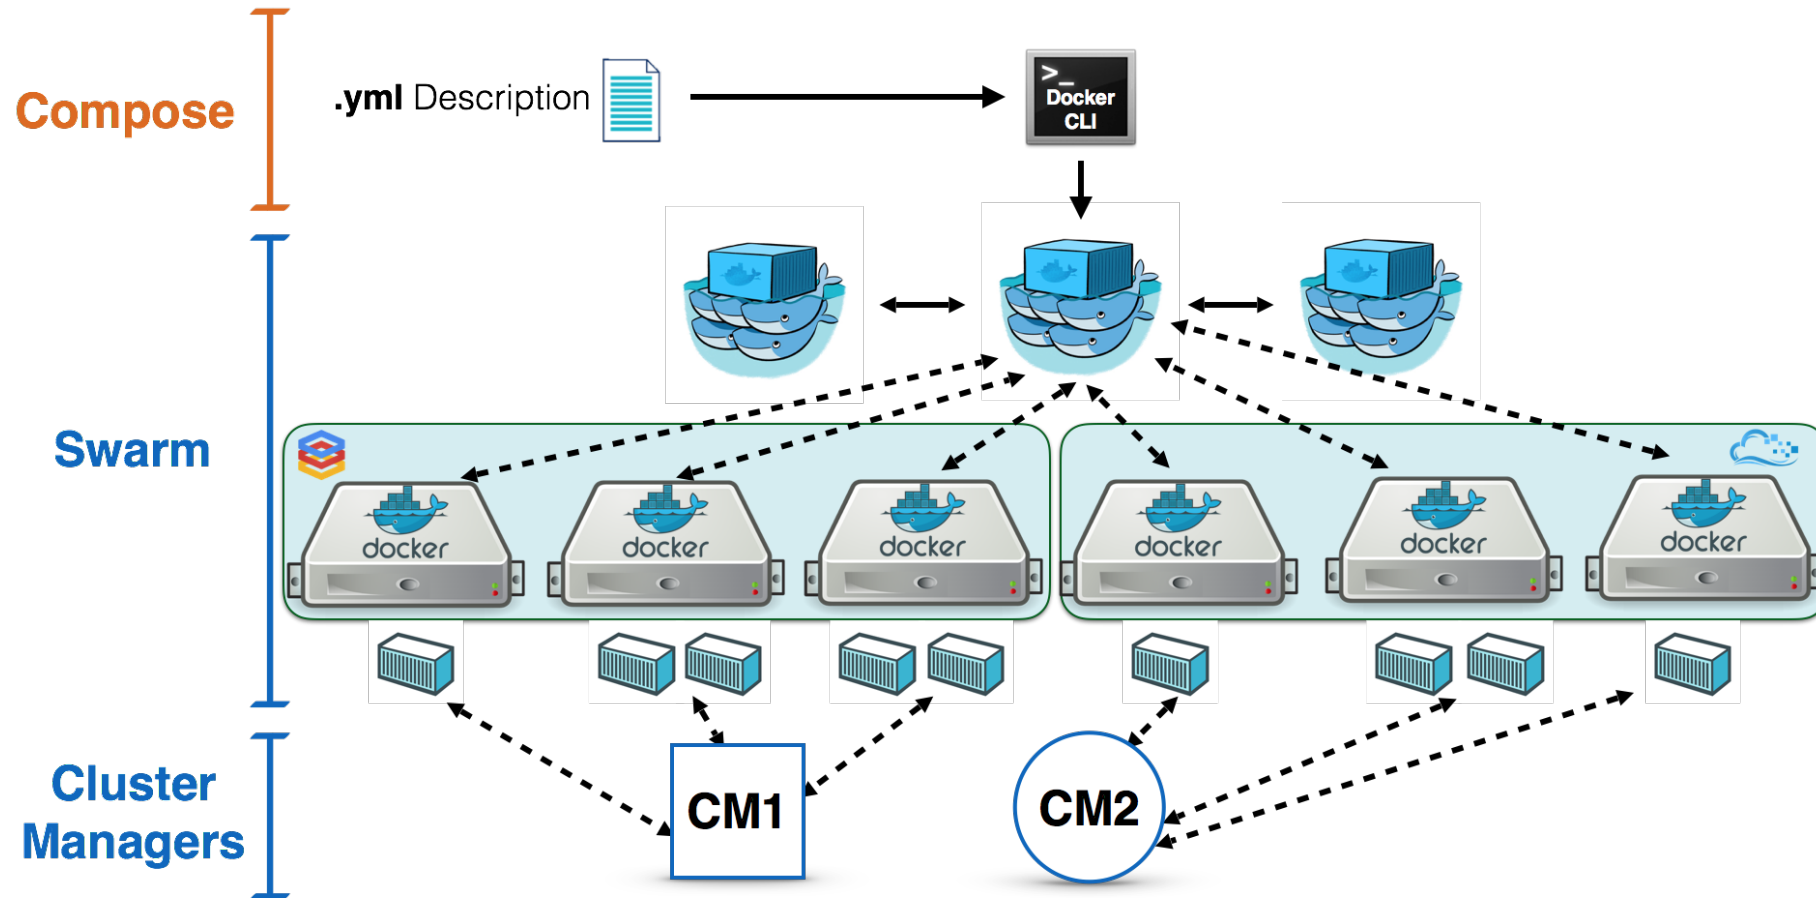
Recap

Recap Docker

- ❑ Containers run on **single** Docker host
- ❑ Containers are **ephemeral**
- ❑ Containers can have external persistence
- ❑ Containers do not contain
- ❑ Operating system **matters**

Swarm Overview

□ Describe swarm in a sentence of two, here.



Initialize Swarm Process

- ❑ Turn single host Docker host into a Multi-host with Docker Swarm Mode.
- ❑ All containers are only **deployed** onto the engine.
- ❑ Swarm Mode turns it into a multi-host **cluster-aware** engine.
- ❑ The first node to initialize the Swarm Mode becomes the **manager**.
- ❑ As new nodes join the cluster, they can **adjust** their roles between managers or workers.
- ❑ You **should** run 3-5 managers in a production environment to ensure high availability.

Create Swarm Mode Cluster

- ❑ Swarm Mode is built **into** the Docker CLI. You can find an overview the possibility commands via:

```
docker swarm --help
```

- ❑ The most important one is how to **initialize** Swarm Mode. Initialization is done via *init*, like this:

```
docker swarm init
```

- ❑ After running the command, the Docker Engine knows how to work with a cluster and becomes the **manager**.
- ❑ The results of an initialization is a **token** used to add additional nodes in a secure fashion.

Swarm Init Output

- Once you initialize Docker Swarm, the token output will look something like this:

```
WSwarm initialized: current node (gpnnuxrri0tzaszblvhdrwqsh) is now a manager.
```

To add a worker to this swarm, run the following command:

```
docker swarm join \
```

```
--token SWMTKN-1-
```

```
1xffnpg61r5wgel28ehuh8o3ip5z2iiwmplbemnksfho5akdqy-  
3zkn8ssrkcnjei34f5rmd4n3q \
```

```
172.17.0.65:2377
```


To add a manager to this swarm, run 'docker swarm join-token manager' and follow the instructions.

Creating a Cluster

- ❑ With Swarm Mode enabled, it is possible to **add** additional nodes and issues commands across all of them.
- ❑ If nodes happen to disappear, for example because of a **crash**, the containers which were running on those hosts will be automatically rescheduled onto other available nodes.
- ❑ The rescheduling ensures you do not lose capacity and provides **high-availability**.
- ❑ On each additional node, you wish to add to the cluster, use the Docker CLI to **join** the existing group.
- ❑ Joining is done by pointing the other host to a **current** manager of the cluster.

Creating a Cluster

- ❑ Docker now uses an **additional** port, 2377, for managing the Swarm.
- ❑ The port should be blocked from public access and only accessed by **trusted** users and nodes.
- ❑ We recommend using VPNs or private networks to **secure** access.



Join Cluster

- ❑ The first task is to **obtain** the token required to add a node to the cluster.
- ❑ For demonstration purposes, we'll ask the manager what the **token** is via *swarm join-token*.

```
token=$(docker -H 172.17.0.65:2345 swarm join-token -q  
worker) && echo $token
```

- ❑ NOTE: In production, this token should be stored securely and only accessible by trusted individuals.

Join Cluster

- ❑ On the second host, **join** the cluster by requesting access via the manager.
- ❑ The **token** is provided as an additional parameter, like this:

```
docker swarm join 172.17.0.65:2377 --token $token
```

Join Cluster

- ❑ By **default**, the manager will automatically accept new nodes being added to the cluster.
- ❑ You can **view** all nodes in the cluster using:


```
docker node ls
```

- ❑ Output:

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER
STATUS5re1humv246m6clujn0zng9jy	docker	Ready		
Activerbavg04vwiijyzzjwdjoc3io2 *	docker	Ready	Active	Leader

Overlay Network

- ❑ Swarm Mode also introduces another **networking** model.
- ❑ In **previous** versions, Docker required the use of an external key-value store, such as Consul, to ensure consistency across the network.
- ❑ The need for consensus and KV has now been incorporated **internally** into Docker and no longer depends on external services.



Overlay Network

- ❑ The improved networking approach follows the **same** syntax as previously.
- ❑ The **overlay** network is used to enable containers on different hosts to communicate.
- ❑ Under the covers, this is a Virtual Extensible LAN (**VXLAN**), designed for large scale cloud based deployments.

Overlay Network

- ❑ The following command will create a new overlay network called **skynet**:

```
docker network create -d overlay skynet
```

- ❑ All containers **registered** to this network can communicate with each other, regardless of which node they are deployed onto.

Deploy Service

- ❑ By default, Docker uses a **spread** replication model for deciding which containers should run on which hosts.
- ❑ The spread approach ensures that containers are deployed across the cluster **evenly**.
- ❑ If one of the nodes are **removed** from the cluster, the containers running are spread across the other available nodes.



Deploy Service

- ❑ The new **Services** concept is used to run containers across the cluster.
- ❑ This is a **higher-level** concept than containers.
- ❑ A service allows you to define how applications should be deployed at **scale**.
- ❑ By updating the service, Docker updates the container required in a **managed** way.

Deployment

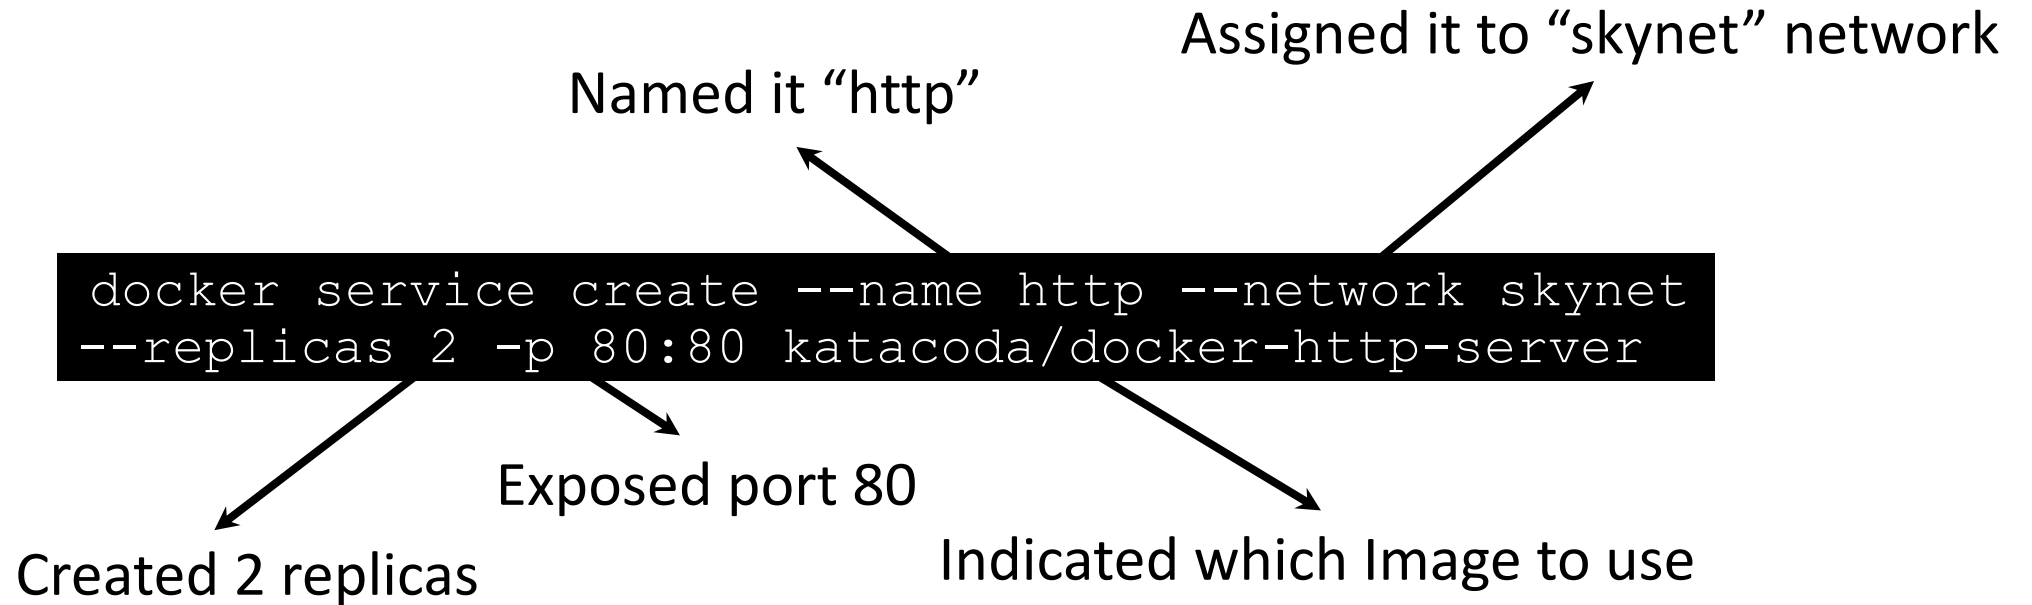
- ❑ As an **example**, let's deploy the Docker Image `katacoda/docker-http-server`.
- ❑ We will give it a friendly **name** like, `http`.
- ❑ And we will **attach** it to the newly created `skynet` network.
- ❑ For ensuring replication and availability, we will run two instances of **replicas**, of the container across our cluster.

Deployment

- ❑ Finally, we will load balance these two containers **together** on port *80*.
- ❑ Sending an HTTP request to any of the nodes in the cluster will process the request by **one** of the containers within the cluster.
- ❑ NOTE: The node which accepts the request might not be the node where the container responds. Instead, Docker load-balances requests across all available containers.
- ❑ So, what should our **code** look like?

Deployment

□ Here it is:



Deployment Verification

- ❑ You can **view** the services running on the cluster using the CLI command

```
docker service ls
```

- ❑ Output:

ID	NAME	MODE	REPLICAS	IMAGE
414yr9hv663w	http	replicated	2/2	katacoda/docker-http-
	server:latest			

Deployment Verification

- ❑ As containers are started you will see them using the `ps` command. You should see one instance of the container on **each** host.
- ❑ **Listing** containers on the first host using `docker ps` will produce:

CONTAINER ID	IMAGE...
<host-1-id>	Image:name

- ❑ Listing containers on the **second** host using `docker ps` produces:

CONTAINER ID	IMAGE...
<host-2-id>	Image:name

Deployment Verification

- ❑ If we issue an HTTP request to the **public** port, it will be processed by the two containers:

```
curl <localhost>
```

- ❑ Output:


```
<h1>This request was processed by host: "host-1-id"</h1>
```

- ❑ If we issue the same command again, what do you think the results will be?

Deployment Verification

- Everything happens just like before, except this request was processed by host 2's container instance:

```
<h1>This request was processed by host: "host-2-id"</h1>
```



Notice which container id is given. Again, using `ps` we could verify this.

What if we do it again?

Health Inspection

- ❑ The Service concept allows you to inspect the health and state of your cluster and the running applications.
- ❑ You can view the list of all the tasks associated with a service across the cluster:

```
docker service ps http
```

- ❑ NOTE: In this case, each task is a container.
- ❑ Output:

ID	NAME	IMAGE	NODE
<service-id-1>	http.1	katacoda/docker-http-server:latest	docker
<service-id-2>	http.2	katacoda/docker-http-server:latest	docker

Health Inspection

- ❑ You can view more details and the configuration of a service by using:

```
docker service inspect --pretty http
```

- ❑ Output:

```
ID:          9b60zdcric52tvo1sskwsta44f
Name:        http
Service Mode: Replicated
Replicas:    2
...
```

Health Inspection

- ❑ On each node, you can ask what tasks it is currently running.

```
docker node ps self
```

- ❑ NOTE: Self refers to the manager node Leader:

- ❑ Output:

ID	NAME	IMAGE	NODE ...
<service-id-1>	http.2	katacoda/docker-http-server:latest	docker ...

Health Inspection

- ❑ Using the ID of a node you can query individual hosts


```
docker node ps $(docker node ls -q | head -n1)
```

- ❑ Output:

ID	NAME	IMAGE	NODE ...
<service-id-2>	http.1	katacoda/docker-http-server:latest	docker ...

Scale Service

- ❑ A Service also allows us to scale how many instances of a task is running across the cluster.
- ❑ As it understands how to launch containers and which containers are running, it can easily start, or remove, containers as required.
- ❑ At the moment the scaling is manual.
 - ❑ However, the API could be hooked up to an external system such as a metrics dashboard.



Scale Service

- ❑ In our example, we had two load-balanced containers running, processing our requests `curl docker`.
- ❑ Command below would scale our `http` service to be run across five containers:

```
docker service scale http=5
```

- ❑ On each host, you will see additional nodes being started.
- ❑ The load balancer will automatically be updated.
- ❑ Requests will now be processed across the new containers.

Result Summary

- The result of this scenario is a two-node Swarm cluster which can run load-balanced containers that can be scaled up and down.

Lab

End of Chapter
